



Level Ancestors

Level Ancestor Queries




Definition: Let v be a vertex at depth d_v in T . For $d \leq d_v$, a *level ancestor query* $\text{LA}(v, d)$ on a vertex v asks to report the ancestor of v at depth d .

Level Ancestor Queries




Definition: Let v be a vertex at depth d_v in T . For $d \leq d_v$, a *level ancestor query* $\text{LA}(v, d)$ on a vertex v asks to report the ancestor of v at depth d .

Level Ancestor Queries




Definition: Let v be a vertex at depth d_v in T . For $d \leq d_v$, a *level ancestor query* $\text{LA}(v, d)$ on a vertex v asks to report the ancestor of v at depth d .

Level Ancestor Queries




Definition: Let v be a vertex at depth d_v in T . For $d \leq d_v$, a *level ancestor query* $\text{LA}(v, d)$ on a vertex v asks to report the ancestor of v at depth d .

Level Ancestor Queries




Definition: Let v be a vertex at depth d_v in T . For $d \leq d_v$, a *level ancestor query* $\text{LA}(v, d)$ on a vertex v asks to report the ancestor of v at depth d .

Level Ancestor Queries



Definition: Let v be a vertex at depth d_v in T . For $d \leq d_v$, a *level ancestor query* $\text{LA}(v, d)$ on a vertex v asks to report the ancestor of v at depth d .

Level Ancestor Queries



Definition: Let v be a vertex at depth d_v in T . For $d \leq d_v$, a *level ancestor query* $\text{LA}(v, d)$ on a vertex v asks to report the ancestor of v at depth d .

The Problem

Given T , design a data structure that is able to preprocess T to answer level ancestors queries.

The Problem

Given T , design a data structure that is able to preprocess T to answer level ancestors queries.

Trivial solutions:

$n = \#$ of nodes

- Preprocessing time: none Size: $O(n)$ Query time: $O(n)$


The Problem

Given T , design a data structure that is able to preprocess T to answer level ancestors queries.

Trivial solutions:

$n = \#$ of nodes

- Preprocessing time: none Size: $O(n)$ Query time: $O(n)$




The Problem

Given T , design a data structure that is able to preprocess T to answer level ancestors queries.

Trivial solutions:

$n = \#$ of nodes

- Preprocessing time: none Size: $O(n)$ Query time: $O(n)$




The Problem

Given T , design a data structure that is able to preprocess T to answer level ancestors queries.

Trivial solutions:

$n = \#$ of nodes

- Preprocessing time: none Size: $O(n)$ Query time: $O(n)$




The Problem

Given T , design a data structure that is able to preprocess T to answer level ancestors queries.

Trivial solutions:

$n = \#$ of nodes

- Preprocessing time: none Size: $O(n)$ Query time: $O(n)$




The Problem

Given T , design a data structure that is able to preprocess T to answer level ancestors queries.

Trivial solutions:

$n = \#$ of nodes

- Preprocessing time: none Size: $O(n)$ Query time: $O(n)$



The Problem

Given T , design a data structure that is able to preprocess T to answer level ancestors queries.

Trivial solutions:

$n = \#$ of nodes

- Preprocessing time: none Size: $O(n)$ Query time: $O(n)$
- Preprocessing time: $O(n^3)$ Size: $O(n^2)$ Query time: $O(1)$
(precompute the answer to all possible queries)

The Problem

Given T , design a data structure that is able to preprocess T to answer level ancestors queries.


Trivial solutions:

$n = \#$ of nodes

- Preprocessing time: none Size: $O(n)$ Query time: $O(n)$
- Preprocessing time: $O(n^3)$ Size: $O(n^2)$ Query time: $O(1)$
- Preprocessing time: $O(n^2)$ Size: $O(n^2)$ Query time: $O(1)$

$$\text{LA}(v, d) = \begin{cases} v & \text{if } d = d_v \\ \text{LA}(\text{parent}(v), d) & \text{if } d < d_v \end{cases}$$


Jump Pointers: Idea



For each vertex v and $\ell = 0, 1, \dots, \lfloor \log d_v \rfloor$, store:

$$J(v, \ell) = \text{LA}(v, d_v - 2^\ell)$$


Jump Pointers: Idea



For each vertex v and $\ell = 0, 1, \dots, \lfloor \log d_v \rfloor$, store:

$$J(v, \ell) = \text{LA}(v, d_v - 2^\ell)$$


Jump Pointers: Idea



For each vertex v and $\ell = 0, 1, \dots, \lfloor \log d_v \rfloor$, store:

$$J(v, \ell) = \text{LA}(v, d_v - 2^\ell)$$


Jump Pointers: Idea



For each vertex v and $\ell = 0, 1, \dots, \lfloor \log d_v \rfloor$, store:

$$J(v, \ell) = \text{LA}(v, d_v - 2^\ell)$$


Jump Pointers: Idea



For each vertex v and $\ell = 0, 1, \dots, \lfloor \log d_v \rfloor$, store:

$$J(v, \ell) = \text{LA}(v, d_v - 2^\ell)$$

Jump Pointers: Idea

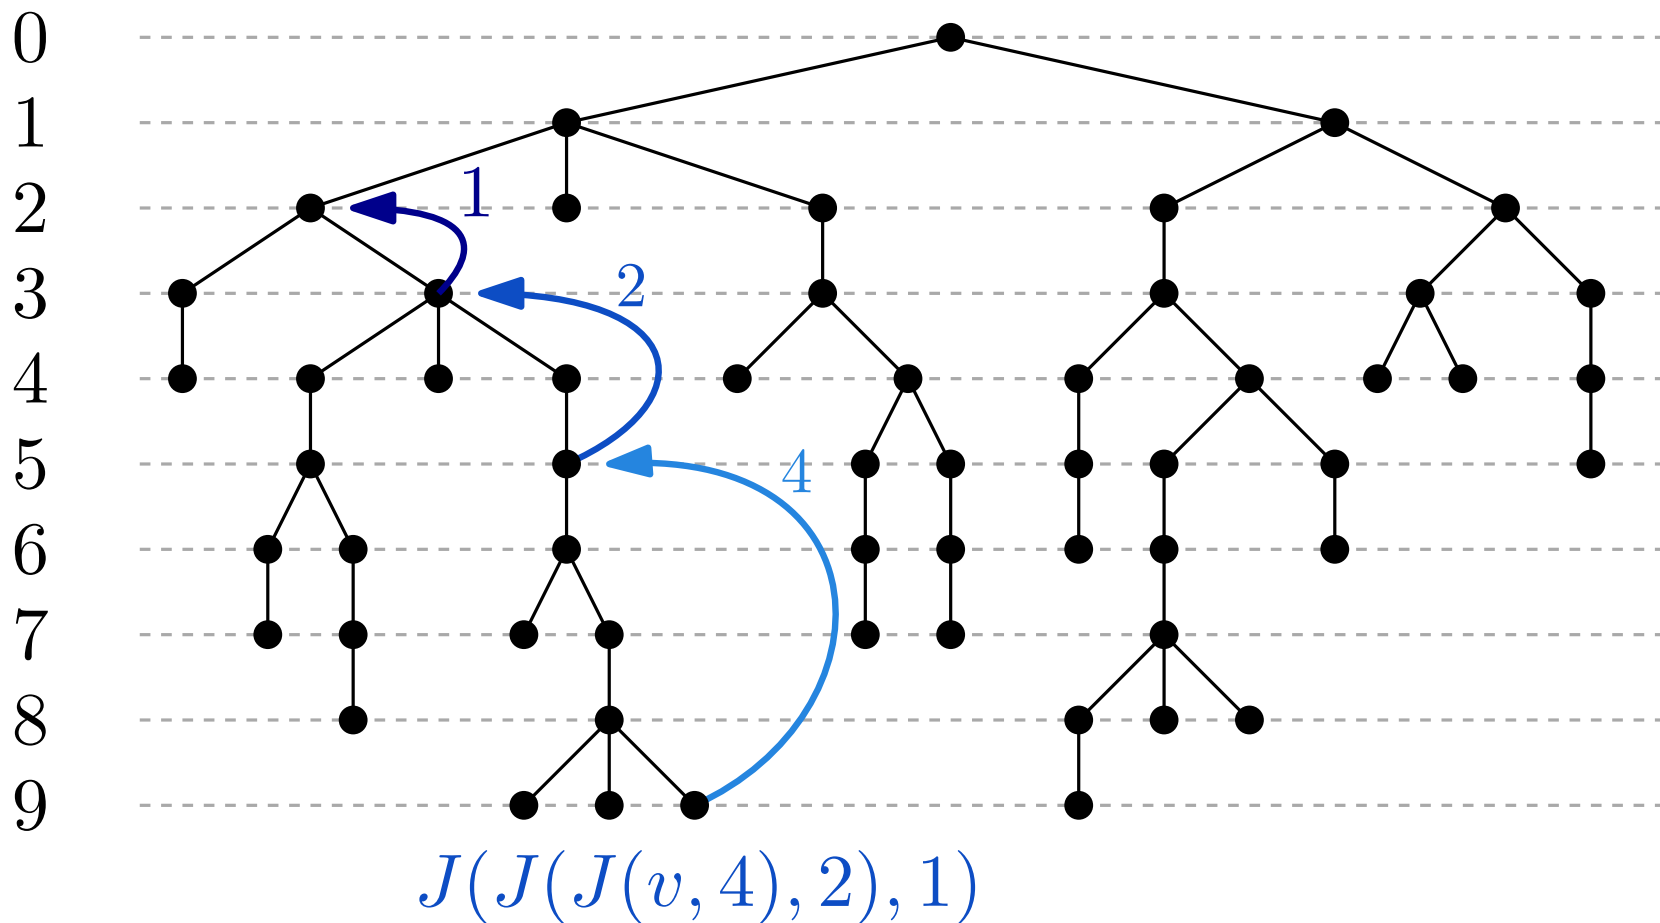


For each vertex v and $\ell = 0, 1, \dots, \lfloor \log d_v \rfloor$, store:

$$J(v, \ell) = \text{LA}(v, d_v - 2^\ell)$$

Total size: $O(n \log h) = O(n \log n)$

Jump Pointers: Query



$$0 < d_v - d = 2^{\ell_k} + 2^{\ell_k-1} + \dots + 2^{\ell_1}$$

$$\ell_{i+1} > \ell_i$$

$$\text{LA}(v, d) = J(\dots J(J(v, \ell_k), \ell_{k-1}), \dots, \ell_1)$$

Number of accessed pointers: $O(\log h) = O(\log n)$

Jump Pointers: Construction

With a DFS visit of T :

- Maintain a stack S that stores all the ancestors of the current vertex v of the visit
- S can be updated in $O(1)$ per traversed edge
- When vertex v is visited, its ancestor at depth d in T is the $(d_v - d)$ -th vertex from the top of the stack

Jump Pointers: Construction

With a DFS visit of T :

- Maintain a stack S that stores all the ancestors of the current vertex v of the visit
- S can be updated in $O(1)$ per traversed edge
- When vertex v is visited, its ancestor at depth d in T is the $(d_v - d)$ -th vertex from the top of the stack

Or using dynamic programming:

$$J(v, \ell) = \begin{cases} \text{parent}(v) & \text{if } \ell = 0 \\ J(J(v, \ell - 1), \ell - 1) & \text{if } \ell > 0 \end{cases}$$

Jump Pointers: Construction

With a DFS visit of T :

- Maintain a stack S that stores all the ancestors of the current vertex v of the visit
- S can be updated in $O(1)$ per traversed edge
- When vertex v is visited, its ancestor at depth d in T is the $(d_v - d)$ -th vertex from the top of the stack

Or using dynamic programming:

$$J(v, \ell) = \begin{cases} \text{parent}(v) & \text{if } \ell = 0 \\ J(J(v, \ell - 1), \ell - 1) & \text{if } \ell > 0 \end{cases}$$

Time complexity: $O(n + n \log h) = O(n \log n)$

Solutions so far

Size	Preprocessing Time	Query Time	Notes
$O(n)$	–	$O(n)$	
$O(n^2)$	$O(n^3)$	$O(1)$	
$O(n^2)$	$O(n^2)$	$O(1)$	

Solutions so far


Size	Preprocessing Time	Query Time	Notes
$O(n)$	–	$O(n)$	
$O(n^2)$	$O(n^3)$	$O(1)$	
$O(n^2)$	$O(n^2)$	$O(1)$	
$O(n \log n)$	$O(n \log n)$	$O(\log n)$	Jump Pointers

Solutions so far

Size	Preprocessing Time	Query Time	Notes
$O(n)$	–	$O(n)$	
$O(n^2)$	$O(n^3)$	$O(1)$	
$O(n^2)$	$O(n^2)$	$O(1)$	
<u>$O(n \log n)$</u>	<u>$O(n \log n)$</u>	<u>$O(\log n)$</u>	Jump Pointers

We want to get rid of the $\log n$ factors!


Long Path Decomposition



Partitions T into a collection of paths \mathcal{D} . Recursively defined:

- Select one of the longest root-to-leaf paths P in T
- Select paths recursively from each the tree of the forest $T \setminus P$


Long Path Decomposition



Partitions T into a collection of paths \mathcal{D} . Recursively defined:

- Select one of the longest root-to-leaf paths P in T
- Select paths recursively from each the tree of the forest $T \setminus P$


Long Path Decomposition



Partitions T into a collection of paths \mathcal{D} . Recursively defined:

- Select one of the longest root-to-leaf paths P in T
- Select paths recursively from each the tree of the forest $T \setminus P$


Long Path Decomposition



Partitions T into a collection of paths \mathcal{D} . Recursively defined:

- Select one of the longest root-to-leaf paths P in T
- Select paths recursively from each the tree of the forest $T \setminus P$


Long Path Decomposition



Partitions T into a collection of paths \mathcal{D} . Recursively defined:

- Select one of the longest root-to-leaf paths P in T
- Select paths recursively from each the tree of the forest $T \setminus P$


Long Path Decomposition



For each path $P_v = \langle v = u_0, \dots, u_k \rangle \in \mathcal{D}$:

- Store an array A_v of length $k + 1$ where $A_v[i]$, $i = 0, \dots, k$, contains (a reference to) u_i
- Each u_i stores a reference $\tau(u_i)$ to v .

Long Path Decomposition




For each path $P_v = \langle v = u_0, \dots, u_k \rangle \in \mathcal{D}$:

- Store an array A_v of length $k + 1$ where $A_v[i]$, $i = 0, \dots, k$, contains (a reference to) u_i
- Each u_i stores a reference $\tau(u_i)$ to v .

$$\text{Total space: } \sum_{P_v \in \mathcal{D}} O(1 + |P_v|) = O(n)$$


Long Path Decomposition



To report $LA(u, d)$:

- Let $v = \tau(u)$
- If $d \geq d_v$: return $A_v[d - d_v]$.


Long Path Decomposition



To report $\text{LA}(u, d)$:

- Let $v = \tau(u)$
- If $d \geq d_v$: return $A_v[d - d_v]$.
- If $d < d_v$: return $\text{LA}(\text{parent}(v), d)$. (recursively)


Long Path Decomposition



To report $\text{LA}(u, d)$:

- Let $v = \tau(u)$
- If $d \geq d_v$: return $A_v[d - d_v]$.
- If $d < d_v$: return $\text{LA}(\text{parent}(v), d)$. (recursively)


Long Path Decomposition



To report $\text{LA}(u, d)$:

- Let $v = \tau(u)$
- If $d \geq d_v$: return $A_v[d - d_v]$.
- If $d < d_v$: return $\text{LA}(\text{parent}(v), d)$. (recursively)

Long Path Decomposition



To report $\text{LA}(u, d)$:

- Let $v = \tau(u)$
- If $d \geq d_v$: return $A_v[d - d_v]$.
- If $d < d_v$: return $\text{LA}(\text{parent}(v), d)$. (recursively)

Time: $O(\# \text{recursive calls}) = O(\# \text{paths in } \mathcal{D} \text{ from } v \text{ to the root})$.

Long Path Decomposition


Claim: The number of distinct paths in \mathcal{D} encountered in the path P from v to the root in T is $O(\sqrt{n})$.

Long Path Decomposition

Claim: The number of distinct paths in \mathcal{D} encountered in the path P from v to the root in T is $O(\sqrt{n})$.

Proof: Let $v = v_0, v_1, \dots, v_k$ be the vertices at which a new path of \mathcal{D} is encountered while traversing P .

- Let P_i be the path of \mathcal{D} that contains v_i .




Long Path Decomposition

Claim: The number of distinct paths in \mathcal{D} encountered in the path P from v to the root in T is $O(\sqrt{n})$.

Proof: Let $v = v_0, v_1, \dots, v_k$ be the vertices at which a new path of \mathcal{D} is encountered while traversing P .

- Let P_i be the path of \mathcal{D} that contains v_i .
- Let $h(v)$ be the height v in T (i.e., the length of the longest path from v to a leaf of T).




Long Path Decomposition

Claim: The number of distinct paths in \mathcal{D} encountered in the path P from v to the root in T is $O(\sqrt{n})$.

Proof: Let $v = v_0, v_1, \dots, v_k$ be the vertices at which a new path of \mathcal{D} is encountered while traversing P .

- Let P_i be the path of \mathcal{D} that contains v_i .
- Let $h(v)$ be the height v in T (i.e., the length of the longest path from v to a leaf of T).
- By the long-path decomposition, $|P_i| \geq h(v_i) \geq i$.




Long Path Decomposition

Claim: The number of distinct paths in \mathcal{D} encountered in the path P from v to the root in T is $O(\sqrt{n})$.

Proof: Let $v = v_0, v_1, \dots, v_k$ be the vertices at which a new path of \mathcal{D} is encountered while traversing P .

- Let P_i be the path of \mathcal{D} that contains v_i .
- Let $h(v)$ be the height v in T (i.e., the length of the longest path from v to a leaf of T).
- By the long-path decomposition, $|P_i| \geq h(v_i) \geq i$.



$$n \geq \left| \bigcup_{i=1}^k P_i \right| \geq \sum_{i=1}^k i \geq \frac{k^2}{2} \implies \sqrt{2n} \geq k.$$


□

Long Path Decomposition

Claim: The number of distinct paths in \mathcal{D} encountered in the path P from v to the root in T is $O(\sqrt{n})$.

Proof: Let $v = v_0, v_1, \dots, v_k$ be the vertices at which a new path of \mathcal{D} is encountered while traversing P .

- Let P_i be the path of \mathcal{D} that contains v_i .
- Let $h(v)$ be the height v in T (i.e., the length of the longest path from v to a leaf of T).
- By the long-path decomposition, $|P_i| \geq h(v_i) \geq i$.




$$n \geq \left| \bigcup_{i=1}^k P_i \right| \geq \sum_{i=1}^k i \geq \frac{k^2}{2} \implies \sqrt{2n} \geq k.$$

□

Time: $O(\sqrt{n})$

Is this tight?

Long Path Decomposition



Time: $\Omega(\sqrt{n})$


Solutions so far

Size	Preprocessing Time	Query Time	Notes
$O(n)$	–	$O(n)$	
$O(n^2)$	$O(n^3)$	$O(1)$	
$O(n^2)$	$O(n^2)$	$O(1)$	
$O(n \log n)$	$O(n \log n)$	$O(\log n)$	Jump Pointers

Solutions so far

Size	Preprocessing Time	Query Time	Notes
$O(n)$	–	$O(n)$	
$O(n^2)$	$O(n^3)$	$O(1)$	
$O(n^2)$	$O(n^2)$	$O(1)$	
$O(n \log n)$	$O(n \log n)$	$O(\log n)$	Jump Pointers
$O(n)$	$O(n)$	$O(\sqrt{n})$	Long Path Dec.


Long Path Decomposition + Ladders



Let $\eta(P_v)$ be the number of vertices of the path $P_v \in \mathcal{D}$.

Extend each path $P_v \in \mathcal{D}$ into a *ladder* L_v with $\eta(P_v)$ more vertices towards the root (if they exist).


Long Path Decomposition + Ladders



Let $\eta(P_v)$ be the number of vertices of the path $P_v \in \mathcal{D}$.

Extend each path $P_v \in \mathcal{D}$ into a *ladder* L_v with $\eta(P_v)$ more vertices towards the root (if they exist).


Long Path Decomposition + Ladders



Let $\eta(P_v)$ be the number of vertices of the path $P_v \in \mathcal{D}$.

Extend each path $P_v \in \mathcal{D}$ into a *ladder* L_v with $\eta(P_v)$ more vertices towards the root (if they exist).

Long Path Decomposition + Ladders



Let $\eta(P_v)$ be the number of vertices of the path $P_v \in \mathcal{D}$.

Extend each path $P_v \in \mathcal{D}$ into a *ladder* L_v with $\eta(P_v)$ more vertices towards the root (if they exist).

Long Path Decomposition + Ladders

For each ladder $L_v = \langle v' = u_0, u_1, \dots, v = u_j, \dots, u_k \rangle$:

- Store, in v , an array B_v of length $k + 1$ where $B_v[i]$ contains (a reference to) u_i
- Each u_i with $i \geq j$ stores a reference $\tau(u_i)$ to v .

The length of B_v is at most twice the length of $A_v \implies$ the total size is still $O(n)$.

Long Path Decomposition + Ladders

For each ladder $L_v = \langle v' = u_0, u_1, \dots, v = u_j, \dots, u_k \rangle$:


- Store, in v , an array B_v of length $k + 1$ where $B_v[i]$ contains (a reference to) u_i
- Each u_i with $i \geq j$ stores a reference $\tau(u_i)$ to v .

The length of B_v is at most twice the length of $A_v \implies$ the total size is still $O(n)$.

To report $\text{LA}(u, d)$:

- Let $v = \tau(u)$ and $v' = B_v[0]$.
- If $d \geq d_{v'}$: return $B_v[d - d_{v'}]$.
- If $d < d_{v'}$: return $\text{LA}(v', d)$. (recursively)

How many recursive calls?



Long Path Decomposition + Ladders

How many recursive calls?


- If we recurse, L_v cannot contain the root of T .

$$|L_v| = \eta(L_v) - 1 = 2\eta(P_v) - 1$$

- Since $u \in P_v$ we have:

$$h(v') \geq |L_v| = 2\eta(P_v) - 1 \geq 2(1 + h(u)) - 1 \geq 2h(u) + 1$$

- The height of the queried vertex doubles at every iteration
 $\implies O(\log n)$ iterations.




Solutions so far

Size	Preprocessing Time	Query Time	Notes
$O(n)$	–	$O(n)$	
$O(n^2)$	$O(n^3)$	$O(1)$	
$O(n^2)$	$O(n^2)$	$O(1)$	
$O(n \log n)$	$O(n \log n)$	$O(\log n)$	Jump Pointers
$O(n)$	$O(n)$	$O(\sqrt{n})$	Long Path Dec.

Solutions so far

Size	Preprocessing Time	Query Time	Notes
$O(n)$	–	$O(n)$	
$O(n^2)$	$O(n^3)$	$O(1)$	
$O(n^2)$	$O(n^2)$	$O(1)$	
$O(n \log n)$	$O(n \log n)$	$O(\log n)$	Jump Pointers
$O(n)$	$O(n)$	$O(\sqrt{n})$	Long Path Dec.
$O(n)$	$O(n)$	$O(\log n)$	+ Ladders

Long Path Dec. + Ladders + Jump Pointers



$$0 < d_u - d = 2^{\ell_k} + 2^{\ell_k-1} + \dots + 2^{\ell_1}$$

To report $\text{LA}(u, d)$:

- Let $w = J(u, \ell_k)$, $v = \tau(w)$ and $v' = B_v[0]$.
- Return $B_v[d_{v'} - d]$.

Query time: $O(1)$

Long Path Dec. + Ladders + Jump Pointers

Jump Pointers

Space usage: $O(\underbrace{n}_{\text{Ladders}} + \overbrace{n \log n}^{\text{Jump Pointers}}) = O(n \log n)$

Long Path Dec. + Ladders + Jump Pointers

$$\text{Space usage: } O(\underbrace{n}_{\text{Ladders}} + \underbrace{n \log n}_{\text{Jump Pointers}}) = O(n \log n)$$

A trick to reduce space:

- Only store jump pointers $J(v, \ell)$ in the leaves v of T .
- For each node u of T , store a reference to a leaf λ_u in the subtree of T rooted at u .
- $\text{LA}(u, d) = \text{LA}(\lambda_u, d)$

Long Path Dec. + Ladders + Jump Pointers

$$\text{Space usage: } O(\underbrace{n}_{\text{Ladders}} + \underbrace{n \log n}_{\text{Jump Pointers}}) = O(n \log n)$$

A trick to reduce space:

- Only store jump pointers $J(v, \ell)$ in the leaves v of T .
- For each node u of T , store a reference to a leaf λ_u in the subtree of T rooted at u .
- $\text{LA}(u, d) = \text{LA}(\lambda_u, d)$

Space usage: $O(n + L \log n)$, where $L = \# \text{leaves of } T$.

Solutions so far

Size	Preprocessing Time	Query Time	Notes
$O(n)$	–	$O(n)$	
$O(n^2)$	$O(n^3)$	$O(1)$	
$O(n^2)$	$O(n^2)$	$O(1)$	
$O(n \log n)$	$O(n \log n)$	$O(\log n)$	Jump Pointers
$O(n)$	$O(n)$	$O(\sqrt{n})$	Long Path Dec.
$O(n)$	$O(n)$	$O(\log n)$	+ Ladders

Solutions so far


Size	Preprocessing Time	Query Time	Notes
$O(n)$	–	$O(n)$	
$O(n^2)$	$O(n^3)$	$O(1)$	
$O(n^2)$	$O(n^2)$	$O(1)$	
$O(n \log n)$	$O(n \log n)$	$O(\log n)$	Jump Pointers
$O(n)$	$O(n)$	$O(\sqrt{n})$	Long Path Dec.
$O(n)$	$O(n)$	$O(\log n)$	+ Ladders
$O(n + L \log n)$	$O(n + L \log n)$	$O(1)$	+ Ladders, JP

Solutions so far


Size	Preprocessing Time	Query Time	Notes
$O(n)$	–	$O(n)$	
$O(n^2)$	$O(n^3)$	$O(1)$	
$O(n^2)$	$O(n^2)$	$O(1)$	
$O(n \log n)$	$O(n \log n)$	$O(\log n)$	Jump Pointers
$O(n)$	$O(n)$	$O(\sqrt{n})$	Long Path Dec.
$O(n)$	$O(n)$	$O(\log n)$	+ Ladders
$O(n + \underline{L \log n})$	$O(n + \underline{L \log n})$	$O(1)$	+ Ladders, JP

If only we had $O(\frac{n}{\log n})$ leaves...

Macro-Micro trees




Macro-Micro trees



Find the set M of all maximally deep vertices with at least $x = \frac{1}{4} \log n$ descendants.

Macro-Micro trees


Example for $x = 7$



Find the set M of all maximally deep vertices with at least $x = \frac{1}{4} \log n$ descendants.

Macro-Micro trees

Example for $x = 7$




Find the set M of all maximally deep vertices with at least $x = \frac{1}{4} \log n$ descendants.

Split T into a **macro-tree** T' containing all the ancestors of the vertices in M and several **micro-trees** in $T \setminus T'$.

Macro-Micro trees

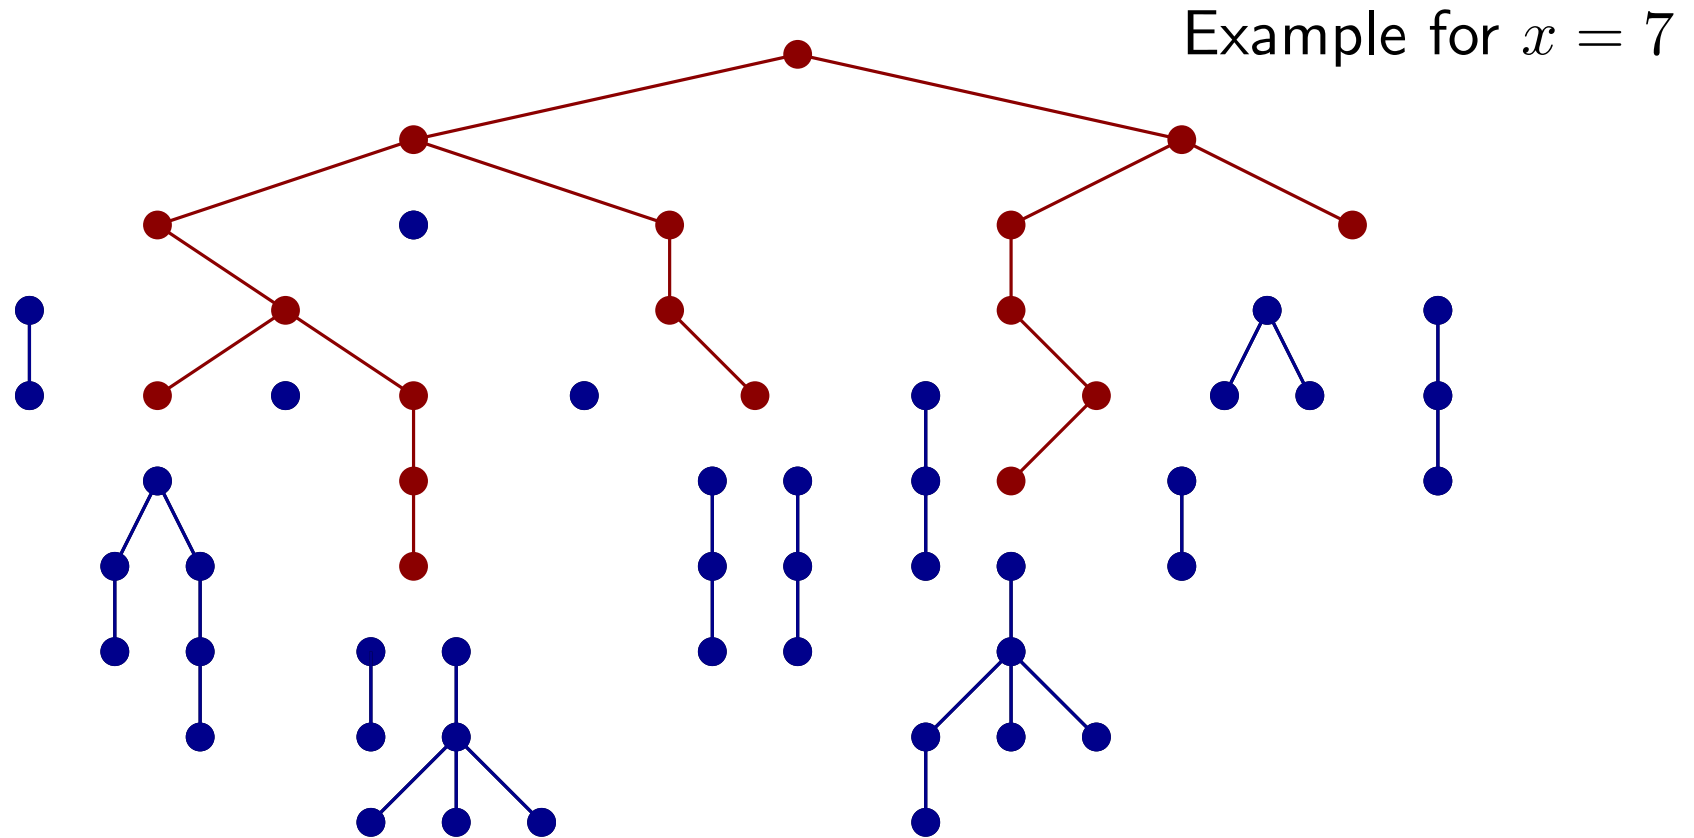
Example for $x = 7$



Find the set M of all maximally deep vertices with at least $x = \frac{1}{4} \log n$ descendants.

Split T into a **macro-tree** T' containing all the ancestors of the vertices in M and several **micro-trees** in $T \setminus T'$.


Macro-Micro trees



Find the set M of all maximally deep vertices with at least $x = \frac{1}{4} \log n$ descendants.

Split T into a **macro-tree** T' containing all the ancestors of the vertices in M and several **micro-trees** in $T \setminus T'$.

Handling the Macro-tree



How many leaves in T' ? The leaves of T' are the vertices in M .


Each vertex in M has at least $\frac{1}{4} \log n$ descendants in T .

$$|M| \cdot \frac{1}{4} \log n \leq n \implies |M| = O\left(\frac{n}{\log n}\right).$$

Build the previous LA oracle \mathcal{O}' on T' .


Size/build time: $O(n + |M| \log n) = O(n)$. **Query time:** $O(1)$.

Handling the Micro-trees




How many different *types* of micro-trees?

A rooted tree on $\leq m$ vertices can be uniquely represented by an array of $2(m - 1)$ bits.





Handling the Micro-trees




How many different *types* of micro-trees?

A rooted tree on $\leq m$ vertices can be uniquely represented by an array of $2(m - 1)$ bits.

Perform a DFS traversal. Write 0 when an edge is traversed towards the leaves, and 1 when it is traversed towards the root.



Handling the Micro-trees



How many different *types* of micro-trees?

A rooted tree on $\leq m$ vertices can be uniquely represented by an array of $2(m - 1)$ bits.

Perform a DFS traversal. Write 0 when an edge is traversed towards the leaves, and 1 when it is traversed towards the root.





Diagram illustrating a 1D lattice structure with sites indexed from 1 to $2(m-1)$. The first two sites are labeled 1 and 2, and the last site is labeled $2(m-1)$. The sites are represented by a horizontal row of boxes, with the first two boxes containing the value 0.


Handling the Micro-trees




How many different *types* of micro-trees?

A rooted tree on $\leq m$ vertices can be uniquely represented by an array of $2(m - 1)$ bits.

Perform a DFS traversal. Write 0 when an edge is traversed towards the leaves, and 1 when it is traversed towards the root.

[illegible]


Handling the Micro-trees



How many different *types* of micro-trees?


A rooted tree on $\leq m$ vertices can be uniquely represented by an array of $2(m - 1)$ bits.

Perform a DFS traversal. Write 0 when an edge is traversed towards the leaves, and 1 when it is traversed towards the root.



1	2	...										$2(m - 1)$
0	0	1	1	0								


Handling the Micro-trees



How many different *types* of micro-trees?


A rooted tree on $\leq m$ vertices can be uniquely represented by an array of $2(m - 1)$ bits.

Perform a DFS traversal. Write 0 when an edge is traversed towards the leaves, and 1 when it is traversed towards the root.



1	2	...										$2(m - 1)$
0	0	1	1	0	0							


Handling the Micro-trees



How many different *types* of micro-trees?


A rooted tree on $\leq m$ vertices can be uniquely represented by an array of $2(m - 1)$ bits.

Perform a DFS traversal. Write 0 when an edge is traversed towards the leaves, and 1 when it is traversed towards the root.



1	2	...										$2(m - 1)$
0	0	1	1	0	0	1						


Handling the Micro-trees



How many different *types* of micro-trees?


A rooted tree on $\leq m$ vertices can be uniquely represented by an array of $2(m - 1)$ bits.

Perform a DFS traversal. Write 0 when an edge is traversed towards the leaves, and 1 when it is traversed towards the root.



1	2	...										$2(m - 1)$
0	0	1	1	0	0	1	0					


Handling the Micro-trees



How many different *types* of micro-trees?


A rooted tree on $\leq m$ vertices can be uniquely represented by an array of $2(m - 1)$ bits.

Perform a DFS traversal. Write 0 when an edge is traversed towards the leaves, and 1 when it is traversed towards the root.



1	2	...											$2(m - 1)$
0	0	1	1	0	0	1	0	1					


Handling the Micro-trees




How many different *types* of micro-trees?

A rooted tree on $\leq m$ vertices can be uniquely represented by an array of $2(m - 1)$ bits.

Perform a DFS traversal. Write 0 when an edge is traversed towards the leaves, and 1 when it is traversed towards the root.



Handling the Micro-trees




How many different *types* of micro-trees?

A rooted tree on $\leq m$ vertices can be uniquely represented by an array of $2(m - 1)$ bits.


Perform a DFS traversal. Write 0 when an edge is traversed towards the leaves, and 1 when it is traversed towards the root.

Pad with 1s.



1	2	\dots									$2(m-1)$
0	0	1	1	0	0	1	0	1	1	1	1

Handling the Micro-trees




How many different *types* of micro-trees?

A rooted tree on $\leq m$ vertices can be uniquely represented by an array of $2(m-1)$ bits.

At most $2^{2(m-1)} < 2^{2m}$ trees with up to m vertices

$\implies O(2^{2\frac{1}{4} \log n}) = O(\sqrt{n})$ micro-tree types.


Handling the Micro-trees



For each of the $O(\sqrt{n})$ distinct micro-trees types T_i

- Build the trival oracle \mathcal{O}_i with size/preprocessing time $O(|T_i|^2)$ and query time $O(1)$.

Handling the Micro-trees




For each of the $O(\sqrt{n})$ distinct micro-trees types T_i

- Build the trival oracle \mathcal{O}_i with size/preprocessing time $O(|T_i|^2)$ and query time $O(1)$.

For each vertex u of T that belongs to a micro tree:

- Store, in u , the index i of the *type* of its micro-tree.
- Store, in u , the vertex $\mu(u)$ in T_i corresponding to u .
- Store, in u , the root $\rho(u)$ of its micro-tree.

Handling the Micro-trees



For each of the $O(\sqrt{n})$ distinct micro-trees types T_i

- Build the trival oracle \mathcal{O}_i with size/preprocessing time $O(|T_i|^2)$ and query time $O(1)$.

For each vertex u of T that belongs to a micro tree:

- Store, in u , the index i of the *type* of its micro-tree.
- Store, in u , the vertex $\mu(u)$ in T_i corresponding to u .
- Store, in u , the root $\rho(u)$ of its micro-tree.

Total size/time: $O(\sqrt{n}) \cdot O(\log^2 n) + O(n) = O(n)$.

Answering a Query

To answer $\text{LA}(u, d)$:

- If u is in the macro tree T' : query \mathcal{O}' for $\text{LA}(u, d)$.
- If u is in a micro-tree T'' :
 - If $d < d_{\rho(u)}$: query \mathcal{O}' for $\text{LA}(\text{parent}(\rho(u)), d)$.
 - Otherwise:
 - Let i be the type of the micro-tree containing u .
 - Query \mathcal{O}_i for $\text{LA}(\mu(u), d - d_{\rho(u)})$.
(and map it back to a vertex in T'')

Query time: $O(1)$.

Solutions so far


Size	Preprocessing Time	Query Time	Notes
$O(n)$	–	$O(n)$	
$O(n^2)$	$O(n^3)$	$O(1)$	
$O(n^2)$	$O(n^2)$	$O(1)$	
$O(n \log n)$	$O(n \log n)$	$O(\log n)$	Jump Pointers
$O(n)$	$O(n)$	$O(\sqrt{n})$	Long Path Dec.
$O(n)$	$O(n)$	$O(\log n)$	+ Ladders
$O(n + L \log n)$	$O(n + L \log n)$	$O(1)$	+ Ladders, JP

Solutions so far

Size	Preprocessing Time	Query Time	Notes
$O(n)$	–	$O(n)$	
$O(n^2)$	$O(n^3)$	$O(1)$	
$O(n^2)$	$O(n^2)$	$O(1)$	
$O(n \log n)$	$O(n \log n)$	$O(\log n)$	Jump Pointers
$O(n)$	$O(n)$	$O(\sqrt{n})$	Long Path Dec.
$O(n)$	$O(n)$	$O(\log n)$	+ Ladders
$O(n + L \log n)$	$O(n + L \log n)$	$O(1)$	+ Ladders, JP
$O(n)$	$O(n)$	$O(1)$	+ Macro-Micro trees

Weighted Level Ancestors


Weighted Level Ancestors



Each edge has a positive weight

Query: Given a vertex v and $d \in \mathbb{N}$, report the deepest ancestor u of v such that the distance from u to v is at least d .


Weighted Level Ancestors



Each edge has a positive weight

Query: Given a vertex v and $d \in \mathbb{N}$, report the deepest ancestor u of v such that the distance from u to v is at least d .


Weighted Level Ancestors



Each edge has a positive weight

Query: Given a vertex v and $d \in \mathbb{N}$, report the deepest ancestor u of v such that the distance from u to v is at least d .

Weighted Level Ancestors



Each edge has a positive weight

Query: Given a vertex v and $d \in \mathbb{N}$, report the deepest ancestor u of v such that the distance from u to v is at least d .

Can be solved with $O(n)$ preprocessing, $O(n)$ space, and $O(\log \log D)$ time

[Amit et. al., Dynamic Text and Static Pattern Matching]